

## CyberX – Mind4Future 2024

# Software Exploitation

**Walter Tiberti, Ph.D**

RTD/a in IINF/03-Telecomunicazioni  
*Dipartimento di Ingegneria,  
Scienze dell'Informazione e Matematica*  
Università degli Studi dell'Aquila



**DISIM**

Email: [walter.tiberti@univaq.it](mailto:walter.tiberti@univaq.it)



# Agenda

- ▶ Introduction
- ▶ Buffer Overflow Vulnerabilities
  - Stack structure (review)
  - Simple case
  - Countermeasures
  - Ret2LibC
  - Return-oriented Programming (ROP)
- ▶ Format String Vulnerabilities
- ▶ Heap Exploitation (notions)
- ▶ Exploitation on Windows (notions)
- ▶ Exercises

# Introduction

- ▶ In this lecture we will discuss the most common **memory corruption vulnerabilities**
  - *Buffer overflows*
  - *Format String vuln.*
  - *Heap overflows*
- ▶ We assume we are working on a **x86 64-bit CPU (x86\_64)** and a **Linux OS**
  - I will try to give you some pointers on what changes in different OS/platforms
  - Feel free to ask for more info!

# Buffer Overflow Vulnerabilities

# Buffer Overflow

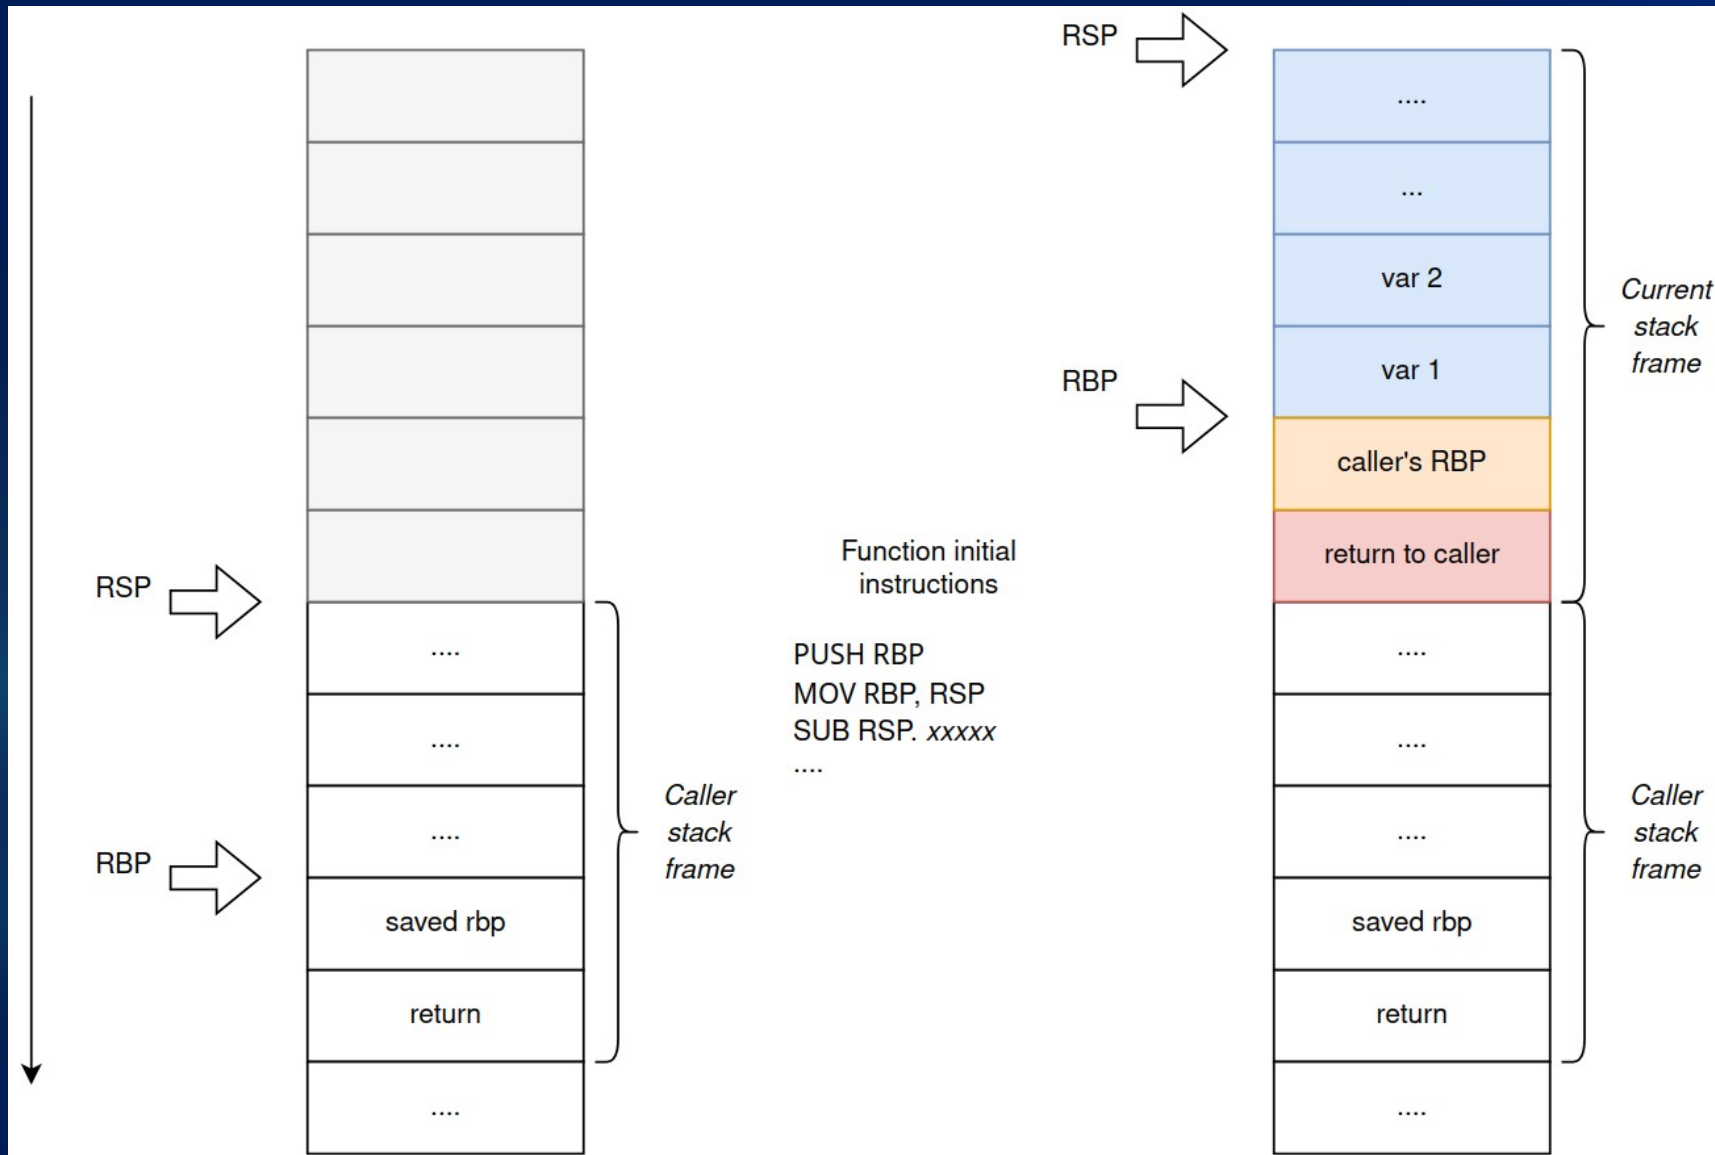
- ▶ **What:** a vulnerability happening when an area of memory (“*buffer*”) is written over exceeding its size – **overwriting what follows**
  - If the buffer is located on the stack → **Stack Buffer overflow**
- ▶ **Causes:** no strong **bound checking**, off-by-1 errors (**NUL** byte), integer over/underflows
- ▶ **Consequences:** depends on where the buffer is allocated, in general:
  - Arbitrary (over)write (sometimes *Write-What-Where*)
  - Code Execution



# CPU Stack review in two slides

- ▶ What: a memory area used as temporary storage, reserved at thread-level, managed by CPU
  - Normally addressed via special purpose registers.  
In case of `x86_64`:
    - `RSP` → 64-bit CPU register used as stack pointer, i.e., always points to the current top of the stack
    - `RBP` → 64-bit CPU register used as frame pointer, i.e., `auxiliary` pointer used to locate the portion of the stack reserved for current function arguments/locals
- ▶ **Last-in-First-Out policy** with `two operations`/instruction:
  - `PUSH x` → reserve space by decrementing the stack pointer and store `x` (word-sized)
  - `POP y` → retrieve the element on the top of the stack and store it in `y`, incrementing the stack pointer
- ▶ Organized in `stack frames` → created at function start, discarded at function end.
- ▶ On `x86_64` Linux ABI, the stack frames contain also:
  - The **frame pointer value before the call** (“old” `RBP` value before the function call)
  - Return address to the caller’s code (!!!)

# CPU Stack review in two slides - example



Remember that:

- The **CALL** instruction automatically push the next address (after the call) onto the stack
- The **RET** instruction automatically retrieve the top of the stack (via RSP) and replace the instruction pointer (RIP) with it

# Stack Buffer overflow

- ▶ Typical issue when there is an array declared as non-static function-local variable and no bound checking is adopted
  - Most compilers/languages allocate thread-local non-static local variable on the stack
  - Security issue: the stack contains also metadata related to code execution
    - Return addresses!
- ▶ Typical attack objective: replace the return address with an address pointing to malicious code (e.g., spawn a shell)



# Vulnerable code – source code view

```
1 #include <stdio.h>
2
3 int main(int argc, char *argv[])
4 {
5     char vulnerable[32]; // what if we write >32 characters?
6     char victim[16]; // <---
7
8     printf("Insert a string: ");
9     fgets(vulnerable, 256, stdin); // Wrong size!
10
11     printf("Vulnerable: %s\n", vulnerable);
12     printf("Victim: %s\n", victim);
13     return 0;
14 }
```

```
wtiberti@x1c6-hook CyberX $ ./stack-example
Insert a string: aaaa
Vulnerable: aaaa

Victim:
```

```
wtiberti@x1c6-hook CyberX $ ./stack-example
Insert a string: aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa
Vulnerable: aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa
Victim:
Segmentation fault (core dumped) Crash!
```

# Vulnerable code – debugger view

- Just before the `call` to `fgets`:

```
gef> x/32gx $rsp
```

```
0x7fffffffda30: 0x00007fffffffdb98
0x7fffffffda40: 0x0000000000000000
0x7fffffffda50: 0x0000000000000000
0x7fffffffda60: 0x0000000000000000
0x7fffffffda70: 0x00007fffffffdb10
0x7fffffffda80: 0x00007fffffffdac0
0x7fffffffda90: 0x0000000010040040
0x7fffffffdaa0: 0x00007fffffffdb98
0x7fffffffdad0: 0x0000000000000001
0x7fffffffdb00: 0x00007ffff7fd000
0x7fffffffdb10: 0xc84f7ec0fca44cfd
```

```
0x0000000010000000
0x0000000000000000
0x00007ffff7fe53e0
0x00007fffffffdb98
0x00007ffff7dc7e08
0x00007fffffffdb98
0x000000000000401136
0xc84f7ec0ff844cfd
0x0000000000000000
0x000000000000403df0
0xc84f6e87b2da4cfd
```

Victim @ 0x0x7fffffffda40

Vulnerable @ 0x7fffffffda50

Return to libc\_start\_main

# Vulnerable code – debugger view

- ▶ AFTER before the `call` to `fgets`

```
gef> x/32gx $rsp
0x7fffffffda30: 0x00007fffffffdb98
0x7fffffffda40: 0x0000000000000000
0x7fffffffda50: 0x6161616161616161
0x7fffffffda60: 0x6161616161616161
0x7fffffffda70: 0x6161616161616161
0x7fffffffda80: 0x6161616161616161
0x7fffffffda90: 0x6161616161616161
0x7fffffffdaa0: 0x00007fffffff000a
0x7fffffffdad0: 0x0000000000000001
0x7fffffffdd00: 0x00007ffff7fd000
0x7fffffffdd10: 0xc84f7ec0fca44cfd
```

```
0x0000000010000000
0x0000000000000000
0x6161616161616161
0x6161616161616161
0x6161616161616161
0x6161616161616161
0x6161616161616161
0xc84f7ec0ff844cfd
0x0000000000000000
0x000000000000403df0
0xc84f6e87b2da4cfd
```

Victim @ 0x0x7fffffffda40

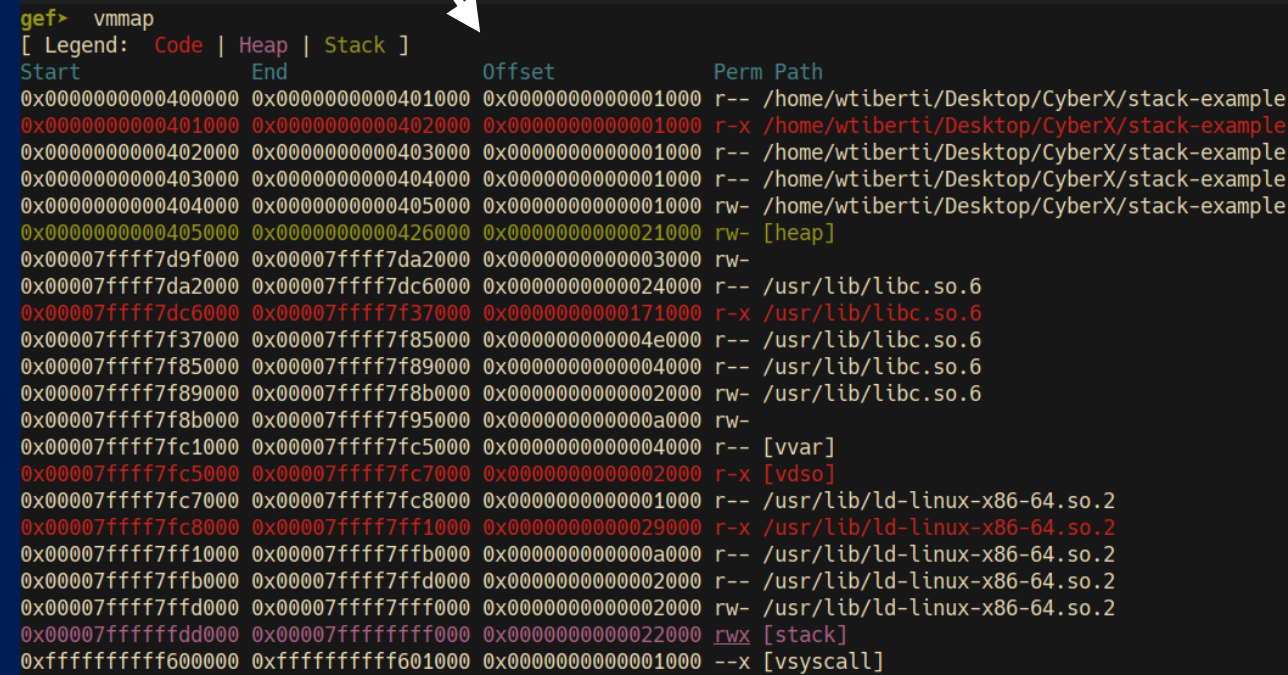
Vulnerable @ 0x7fffffffda50

Return to libc\_start\_main

What happens when `main()` ends?

# After the return

- ▶ We would jump to `0x6161616161616161 ...`
- ▶ That is:
  - **NOT** a *canonical* x86\_64 address
  - **NOT mapped** on the process address space
- ▶ Result → **Crash**  
(segmentation fault)

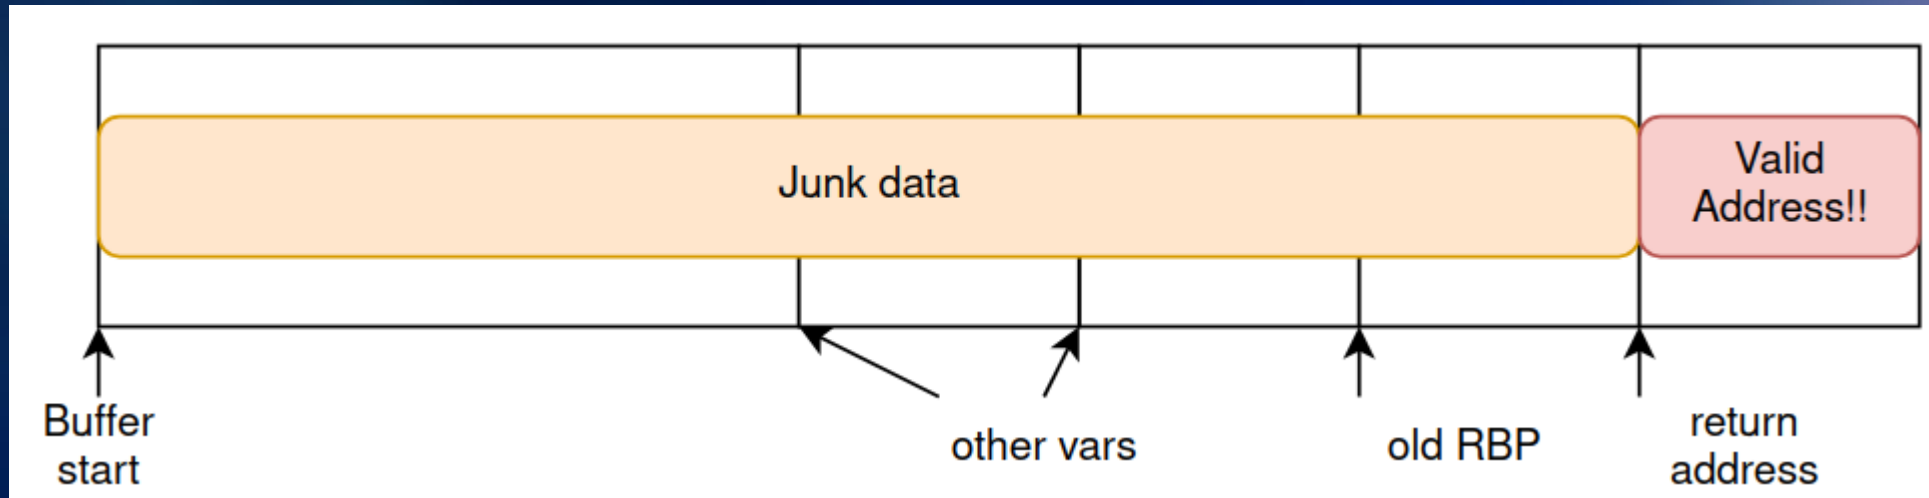


```
gef> vmmap
[ Legend:  Code | Heap | Stack ]
Start      End      Offset   Perm Path
0x0000000000400000 0x0000000000401000 0x0000000000001000 r-- /home/wtiberti/Desktop/CyberX/stack-example
0x0000000000401000 0x0000000000402000 0x0000000000001000 r-x /home/wtiberti/Desktop/CyberX/stack-example
0x0000000000402000 0x0000000000403000 0x0000000000001000 r-- /home/wtiberti/Desktop/CyberX/stack-example
0x0000000000403000 0x0000000000404000 0x0000000000001000 r-- /home/wtiberti/Desktop/CyberX/stack-example
0x0000000000404000 0x0000000000405000 0x0000000000001000 rw- /home/wtiberti/Desktop/CyberX/stack-example
0x0000000000405000 0x0000000000426000 0x0000000000021000 rw- [heap]
0x00007ffff7d9f000 0x00007ffff7da2000 0x0000000000003000 rw-
0x00007ffff7da2000 0x00007ffff7dc6000 0x00000000000024000 r-- /usr/lib/libc.so.6
0x00007ffff7dc6000 0x00007ffff7f37000 0x00000000000171000 r-x /usr/lib/libc.so.6
0x00007ffff7f37000 0x00007ffff7f85000 0x0000000000004e000 r-- /usr/lib/libc.so.6
0x00007ffff7f85000 0x00007ffff7f89000 0x00000000000004000 r-- /usr/lib/libc.so.6
0x00007ffff7f89000 0x00007ffff7f8b000 0x00000000000002000 rw- /usr/lib/libc.so.6
0x00007ffff7f8b000 0x00007ffff7f95000 0x0000000000000a000 rw-
0x00007ffff7f95000 0x00007ffff7fc5000 0x00000000000004000 r-- [vvar]
0x00007ffff7fc5000 0x00007ffff7fc7000 0x00000000000002000 r-x [vdso]
0x00007ffff7fc7000 0x00007ffff7fc8000 0x00000000000001000 r-- /usr/lib/ld-linux-x86-64.so.2
0x00007ffff7fc8000 0x00007ffff7ff1000 0x00000000000029000 r-x /usr/lib/ld-linux-x86-64.so.2
0x00007ffff7ff1000 0x00007ffff7ffb000 0x0000000000000a000 r-- /usr/lib/ld-linux-x86-64.so.2
0x00007ffff7ffb000 0x00007ffff7ffd000 0x00000000000002000 r-- /usr/lib/ld-linux-x86-64.so.2
0x00007ffff7ffd000 0x00007ffff7fff000 0x00000000000002000 rw- /usr/lib/ld-linux-x86-64.so.2
0x00007ffff7fff000 0x00007ffff7fff000 0x000000000000022000 rwx [stack]
0xfffffffff6000000 0xfffffffff6010000 0x00000000000001000 --x [vsyscall]
```



# Attack - idea

- ▶ **Idea:** fill the vulnerable buffer with an input (payload) that contains:
  - The precise amount of characters needed to completely fill the space to reach the interesting metadata (→ the return address)
  - Attacker controlled memory address





# Attack – problems to solve – counting..

- ▶ **First problem:** detect the precise amount of bytes to write for reaching the return address

- Techniques:

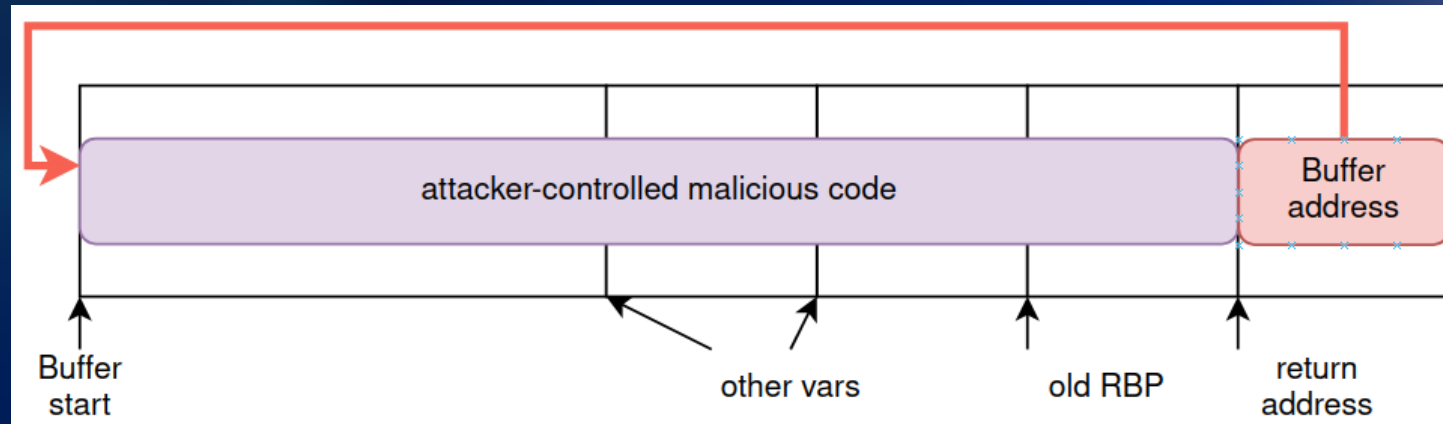
- Manual analysis and counting (time consuming, requires the executable)
- Trial-and-error (time consuming, detectable, black-box)
- Debruijn sequences (faster, require an output)

[https://en.wikipedia.org/wiki/De\\_Bruijn\\_sequence](https://en.wikipedia.org/wiki/De_Bruijn_sequence)

# Attack – problems to solve – where?

- ▶ **Second problem:** what address to use as return address?
- ▶ It depends:
  - The application has interesting code somewhere?
  - The libraries mapped in the process address space have interesting code?
  - No useful code anywhere?
- ▶ **“Old” strategy:** put arbitrary malicious code in the **stack itself** (e.g., in the vuln. buffer) and use the address of it as return address

Shellcode!

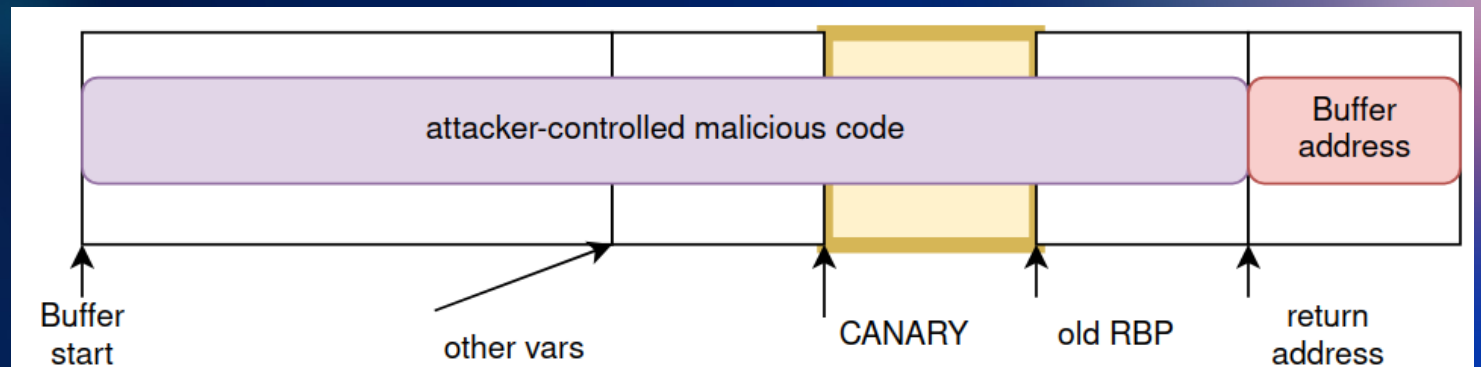


# Countermeasures: DEP/NX

- ▶ The attack strategy has a strong assumption behind: the memory segment containing the stack should be executable
- ▶ The easiest solution is just to avoid assigning executable permissions to the stack!
- ▶ Nowadays this technique is used always and has different names:
  - **DEP** (Data execution prevention)
  - **NX Stack** (Not Executable Stack)
  - **W<sup>X</sup>** (more subtle – *Read XOR Execute* – when a memory segment is writeable, it should not be executable)

# Countermeasures: Stack Canaries (1 of 2)

- ▶ **Idea:** during function prologue, place specific values in the stack before the return address and check them for changes before returning.
  - If there is a Buffer Overflow, the values will change!
  - If so, better crash then giving the attacker a chance to execute code
- ▶ This technique/value is called :
  - Stack Canary
  - Stack Cookie
  - Stack protector



# Countermeasures: Stack Canaries (2 of 2)

- ▶ Types of canaries:
  - **Fixed** (a fixed value decided at compile time)
  - **Terminator** (a value that includes common string terminators e.g., NUL, newline, etc.)
  - **Random** (a random value generated by the OS/kernel)
- ▶ When canaries are deployed, the attacker can still perform the attack if has a way to guess/leak/bruteforce the canary value



# Getting the canary

- ▶ Depending on the scenario, the attacker could
  - Reverse the executable to **get a fixed canary value**
  - Abuse a vulnerable function that **do not stops at terminators** (strcmp vs. memcpy)
  - **Bruteforcing**
- ▶ **Bruteforcing** may work in case of “*forking servers*” application (i.e., an application handling requests by forking and creating child processes) since the **canary is shared** with child processes
  - (just the idea) **what if the attacker just returns to the function itself while providing different canary values? You could get the canary value byte-by-byte**

# PIE and ASLR

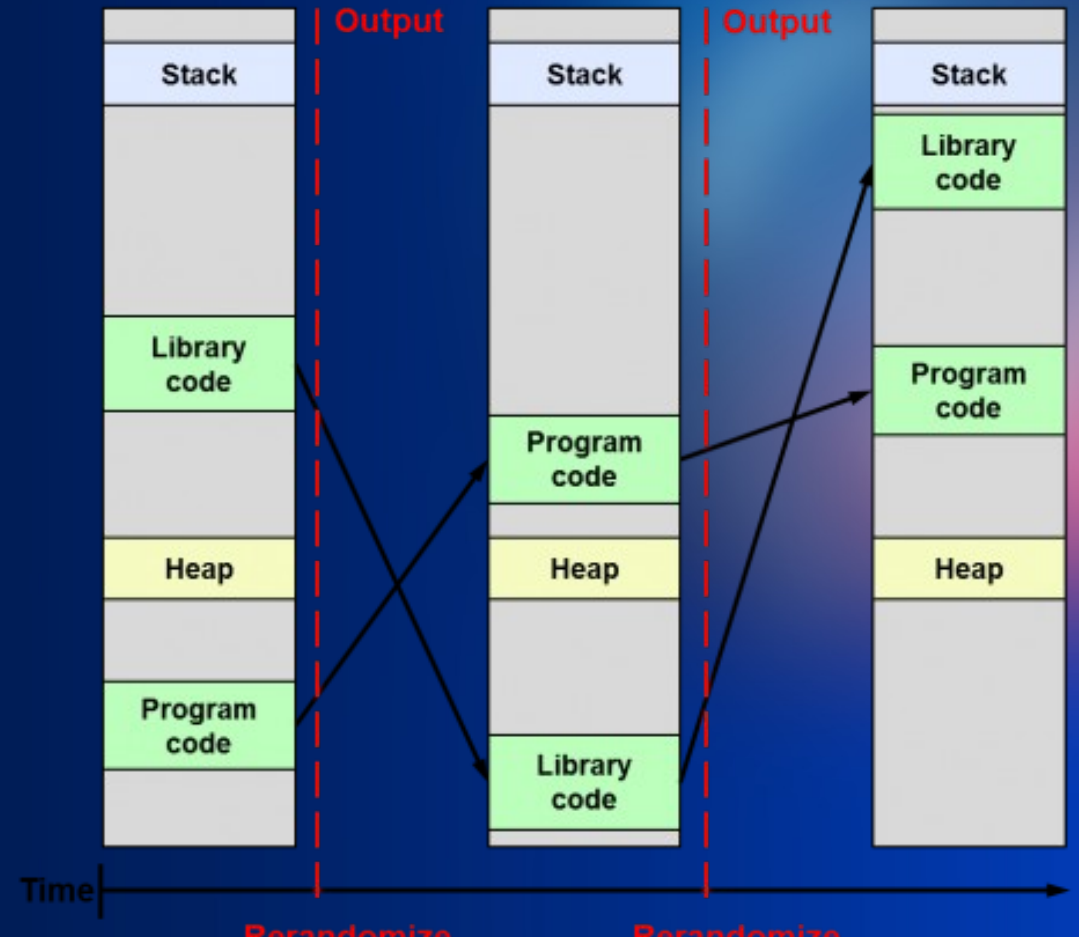
- ▶ As you may already got, knowing which value to put as return address is critical
- ▶ To further block the attackers, a strong defence is to **randomize the position** (i.e., virtual memory addresses) of **memory segments**
- ▶ Two techniques:
  - **PIE** – position independent executable
  - **ASLR** – Address Space Layout Randomization

# PIE

- ▶ PIE is a technique not strictly adopted as security measure but required for building libraries and all those executable that **may not rely on absolute addresses in the code**
- ▶ PIE tells the compiler to **avoid absolute constant addresses** and replace them with an expression like:  
**Base address + Offset**
- ▶ The **base address** depends on where the executable/library is mapped, the **offset** determine the specific location from the start
- ▶ From the security point-of-view, the **offset is known but the base address is not**
- ▶ How to retrieve the base address involve **leaking** addresses in the stack or **bruteforcing**

# ASLR

- ▶ ASLR consists in randomizing the location all the executable segments, so that an attacker has hard time figuring out where things are located
- ▶ Three common variants:
  - No ASLR (0)
  - Just libraries/stack/heap (1)
  - Full ASLR (2)



```
wtiberti@x1c6-hook ~ $ cat /proc/sys/kernel/randomize_va_space
2
```

# Return to LibC


- ▶ In Linux, 90% of the userspace programs uses the C-library. Hence, it is commonly mapped in the process address space
- ▶ The attacker could try to abuse a BoF to call e.g., the C-library function `system("/bin/sh")` to execute a shell. To do so:
  - Retrieve or compute the `system()` address
  - Retrieve or build the string `"/bin/sh"` in a known location (e.g., stack by using RSP)
  - Exploit the BoF putting the address of `"/bin/sh"` in **RDI** (first argument) and the `system()` address as return address



# Return to LibC – system() address

- ▶ In order to compute the address, we need both the **base address** of the C-library when loaded inside the address space of the process and the **offset** to system()
- ▶ The base address **may be trivial or very difficult** to get if ASLR is used – in both cases, leaking/reversing/bruteforcing is used
  - **Note: the libc\_start\_main+X return address is inside the stack!**
  - **Note: ...is there a place in the ELF where imported library function addresses are put?**
- ▶ The **offset** is easy to get (`nm -D /lib/libc.so.6`) **IF** you have the 1:1 C-Library file
  - Different versions, compilation flags etc. may alter the offset
  - There exists databases of compiled C-libraries with symbol address search function: <https://libc.rip/>

# “/bin/sh”

- ▶ **Option 1:** easily found inside libc (`strings -tx /lib/libc.so.6`)
- ▶ **Option 2:** forge the string inside the stack leaving RSP pointing to it
  - `PUSH 0x68732f6e69622f`  `“/bin/sh\x00\x00\x00\x00\x00”`  
in little-endian
- ▶ Using the address forged/found as argument could be not so easy.
- ▶ However, we can use a 2-step approach:
  - 1) Spot the address containing an instruction letting us set RDI with a controlled value. Set this value as return address
  - 2) Manipulate the following stack location (used as return value) so that, after setting RDI, you can return to another location of your choice → `system()`

(more details later)

# Return Oriented Programming (ROP)

- ▶ **Idea**: by exploiting a BoF we can write data starting from the address of a vulnerable buffer. We could **continue writing past the return address** to **forge** small stack frames containing return addresses to **existing fragments of code** ending with the **RET** instruction ("**Gadgets**")
- ▶ By concatenating gadgets (→ creating a "**ROP chain**") we can execute code bypassing DEP/NX entirely
- ▶ This technique is called **Return oriented Programming (ROP)**

# ROP in action – some gadgets

```
mov eax, 0xffffffff
ret
```

```
pop rdi
pop rbp
ret
```

```
pop rbx
pop r12
pop r13
pop rbp
ret
```

```
jmp rax
```

```
sbb eax, eax
ret
```

```
mov qword [rbx], 0
pop rbx
pop r12
pop rbp
ret
```

```
xor eax, eax
call sym.sprintf
```

```
mov eax, ebx
syscall
```

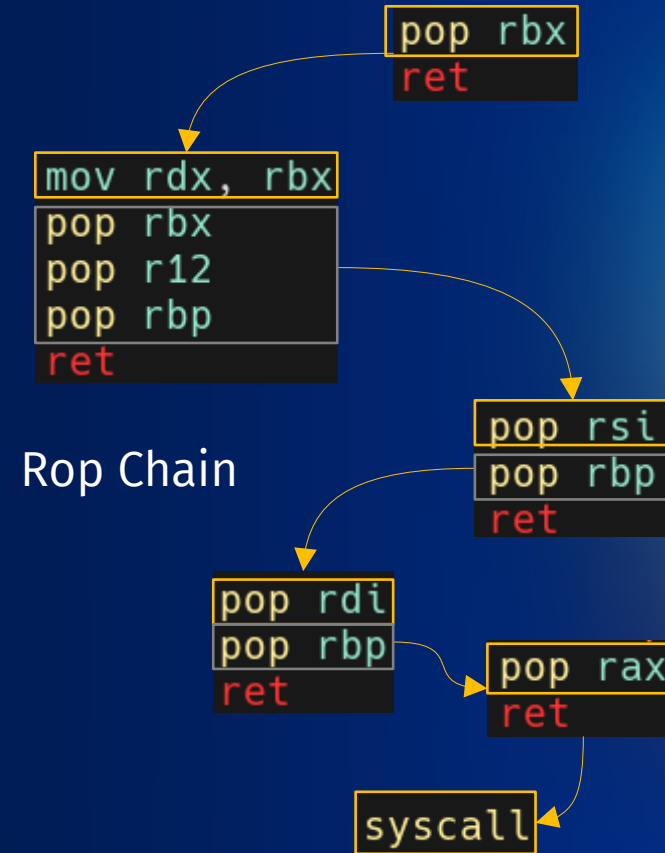
```
mov rbx, qword [rbp - 8]
leave
ret
```

```
syscall
```

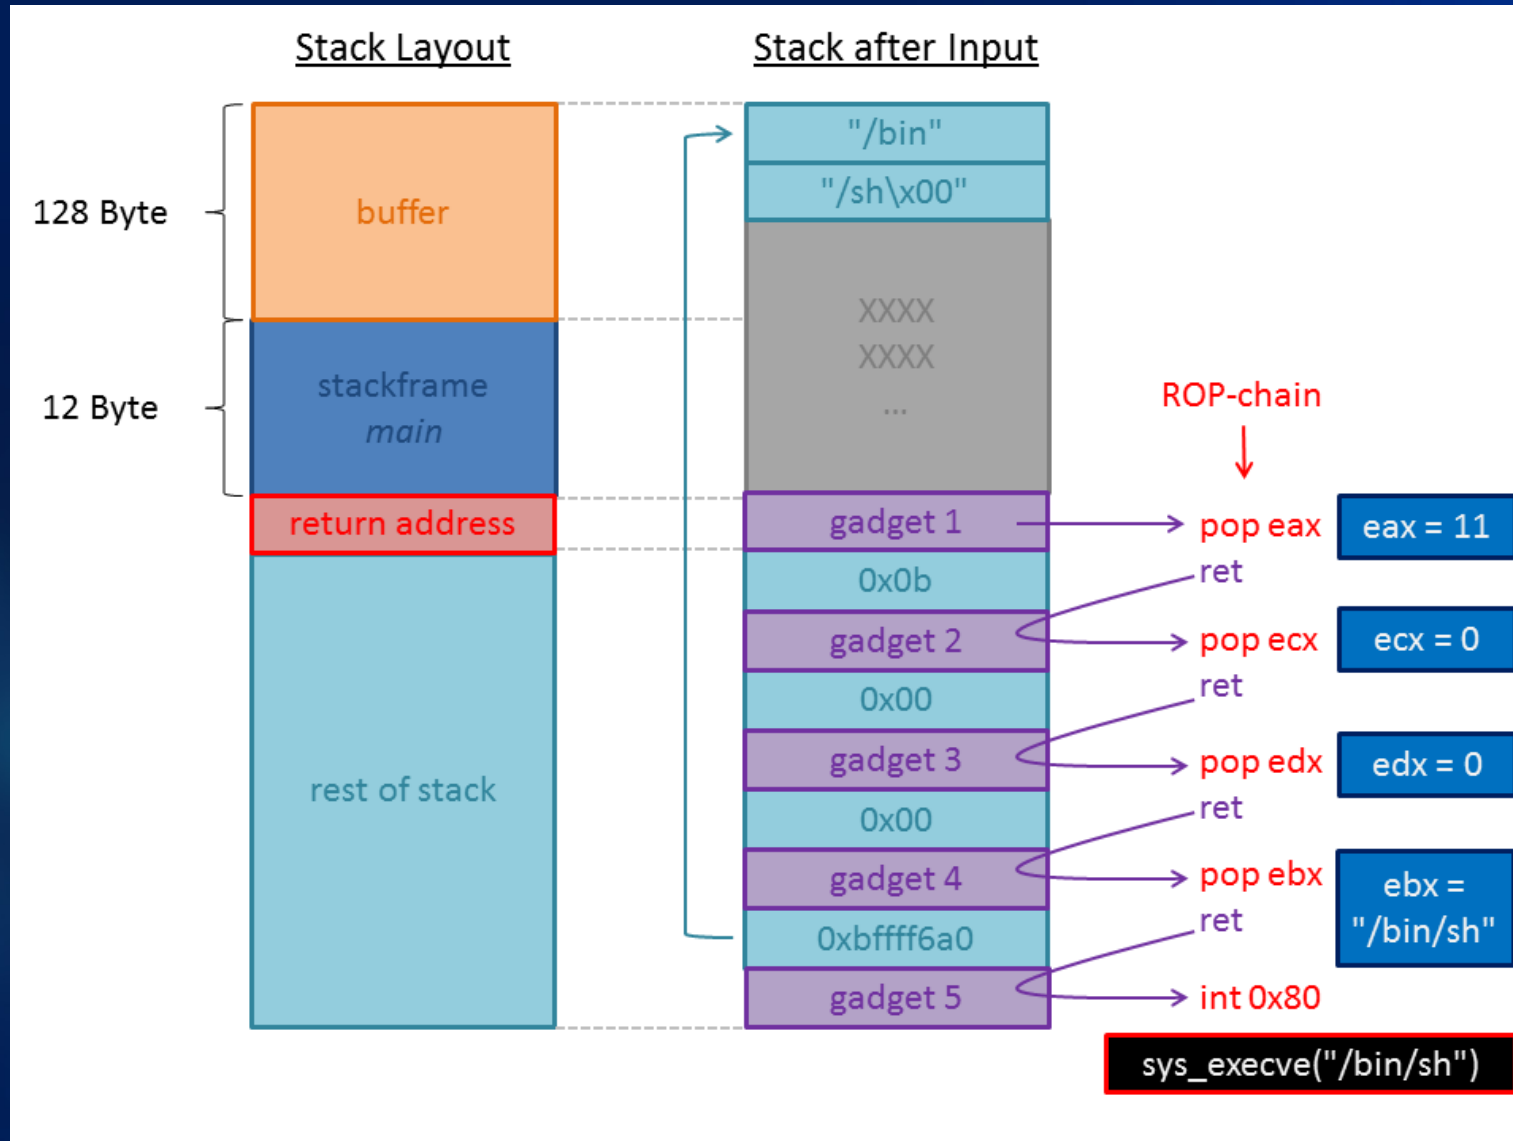
```
call r15
```

```
pop r15
pop rbp
ret
```

```
pop r12
pop r13
pop r14
pop r15
pop rbp
jmp rax
```

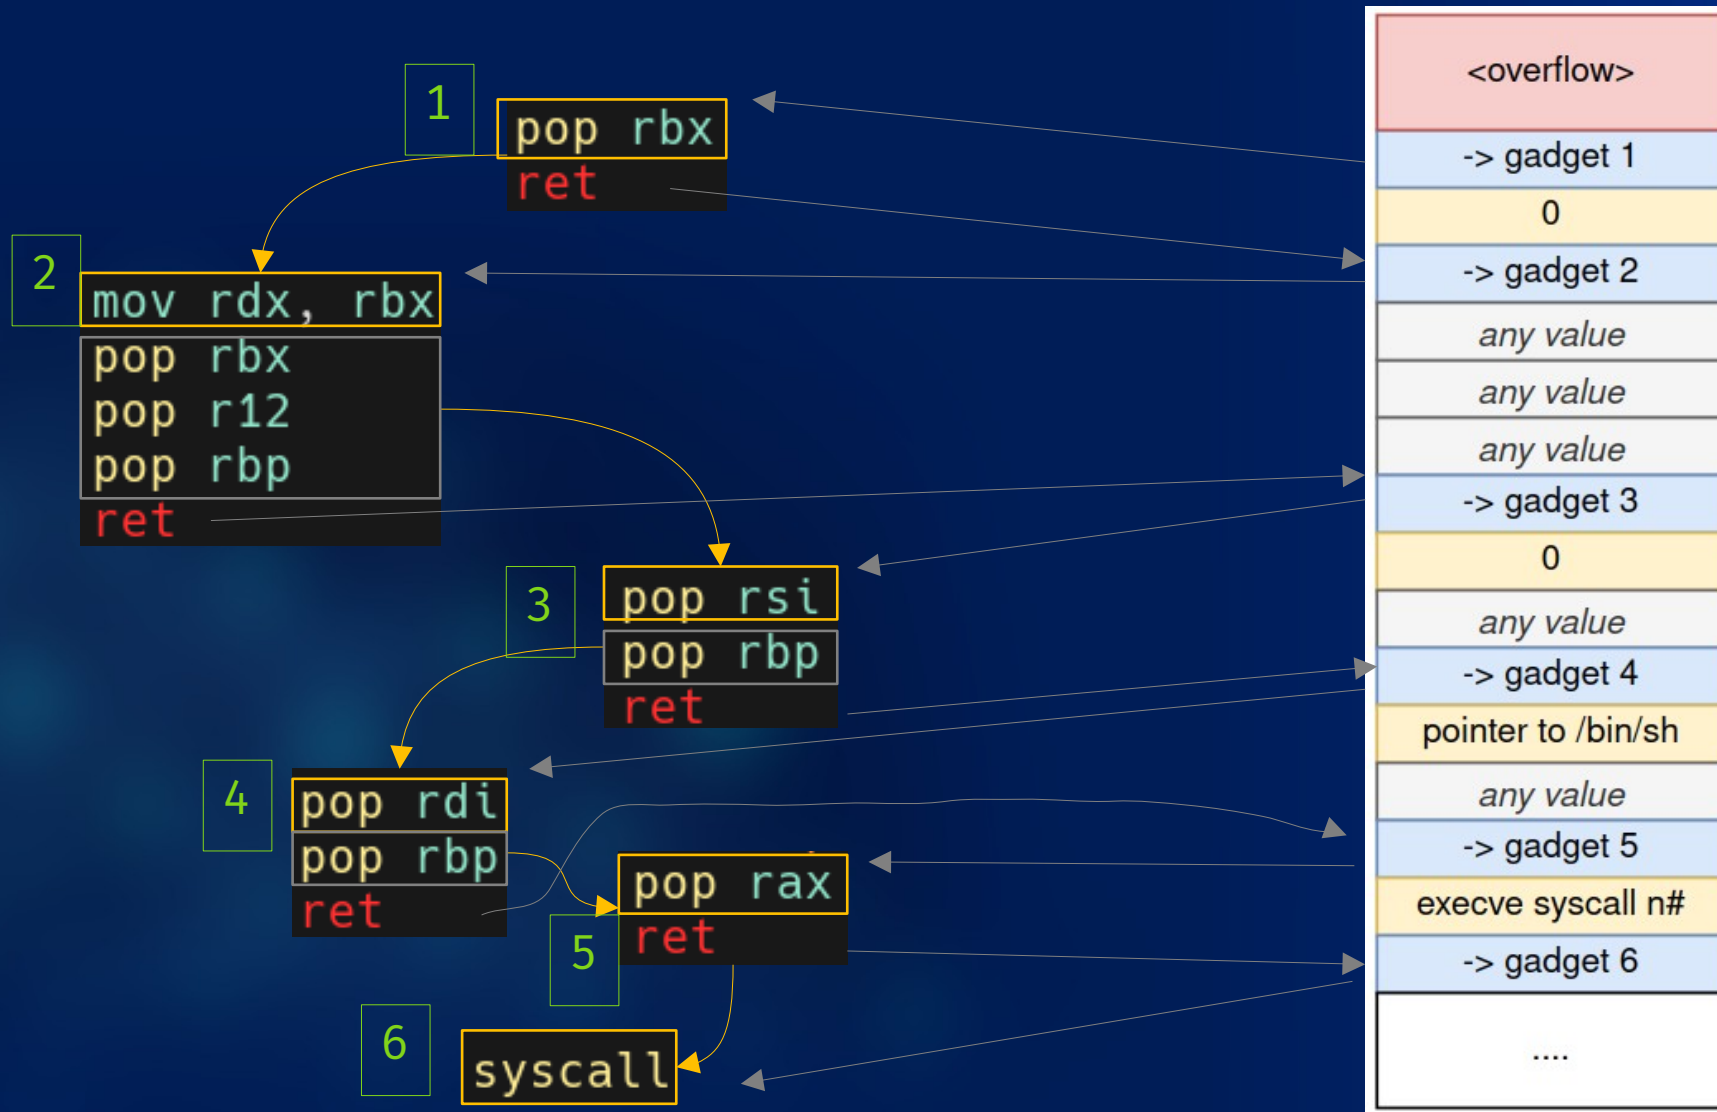


# ROP example (stack view - x86)





# ROP example (stack view - x86\_64)



# How to find gadgets

- ▶ **Option 1:** looking inside the executable and imported libraries (disassembly)
- ▶ **Option 2:** using the **ropper** tool

```
> ropper -f ./a.out
[INFO] Load gadgets for section: LOAD
[LOAD] loading... 100%
[LOAD] removing double gadgets... 100%
```

```
0x00000000000001004: sub rsp, 8; mov rax, qword ptr [rip + 0x2fc1]; test rax, rax; je 0x1016; call rax;
0x0000000000000111a: test byte ptr [rax], al; add byte ptr [rax], al; add byte ptr [rax], al; ret;
0x00000000000001010: test eax, eax; je 0x1016; call rax;
0x00000000000001010: test eax, eax; je 0x1016; call rax; add rsp, 8; ret;
0x0000000000000108b: test eax, eax; je 0x1098; jmp rax;
0x0000000000000108b: test eax, eax; je 0x1098; jmp rax; nop dword ptr [rax]; ret;
0x000000000000010cc: test eax, eax; je 0x10d8; jmp rax;
0x000000000000010cc: test eax, eax; je 0x10d8; jmp rax; nop word ptr [rax + rax]; ret;
0x0000000000000100f: test rax, rax; je 0x1016; call rax;
0x0000000000000100f: test rax, rax; je 0x1016; call rax; add rsp, 8; ret;
0x0000000000000108a: test rax, rax; je 0x1098; jmp rax;
0x0000000000000108a: test rax, rax; je 0x1098; jmp rax; nop dword ptr [rax]; ret;
0x000000000000010cb: test rax, rax; je 0x10d8; jmp rax;
0x000000000000010cb: test rax, rax; je 0x10d8; jmp rax; nop word ptr [rax + rax]; ret;
0x00000000000001080: cld; je 0x1098; mov rax, qword ptr [rip + 0x2f3e]; test rax, rax; je 0x1098; jmp rax;
0x000000000000011a7: cli; sub rsp, 8; add rsp, 8; ret;
0x00000000000001003: cli; sub rsp, 8; mov rax, qword ptr [rip + 0x2fc1]; test rax, rax; je 0x1016; call rax;
0x000000000000011a4: endbr64; sub rsp, 8; add rsp, 8; ret;
0x00000000000001000: endbr64; sub rsp, 8; mov rax, qword ptr [rip + 0x2fc1]; test rax, rax; je 0x1016; call rax;
;
0x0000000000000119f: leave; ret;
0x0000000000000101a: ret;

115 gadgets found
```

# Other topics

- ▶ Other topics omitted due to time constraints:
  - GOT overwrite and RELRO (Partial/Full)
  - Hook rewrite
  - One Gadget
  - Universal Gadget
  - Fini Overwrite

# Exercises

- ▶ For simple BoF the best is to manually craft vulnerable executable and exploit them by adding countermeasures one by one and simulating a remote connection
  - `-zexecstack`, `-fno-stack-protector`, `-no-pie` etc.
  - Use `socat`: `socat TCP-L:12345,fork,reuseaddr EXEC:./example`
- ▶ Protostar/Phoenix: <https://exploit.education/protostar/>
- ▶ ROP Emporium: <https://ropemporium.com/>

# Format String Vulnerabilities



# Format String Vulnerabilities

- ▶ A **format string** is a string consisting of “normal” characters and zero or more “**format string conversion specifiers**”
- ▶ Format strings are largely used in C, Python, PHP, etc. to process, parse and format strings
- ▶ Example:

```
1 #include <stdio.h>
2
3
4 int main(int argc, char *argv[])
5 {
6     int a = 3;
7     unsigned long x = 0xAABBCCDD;
8     char *s1 = "this_is_a_string";
9     double d = 0.0000123;
10
11     printf("a=%d x=%lu s1=%s d=%.6lf \n", a, x, s1, d);
12     return 0;
13 }
```

Format string! (1st arg)

# Format specifiers syntax

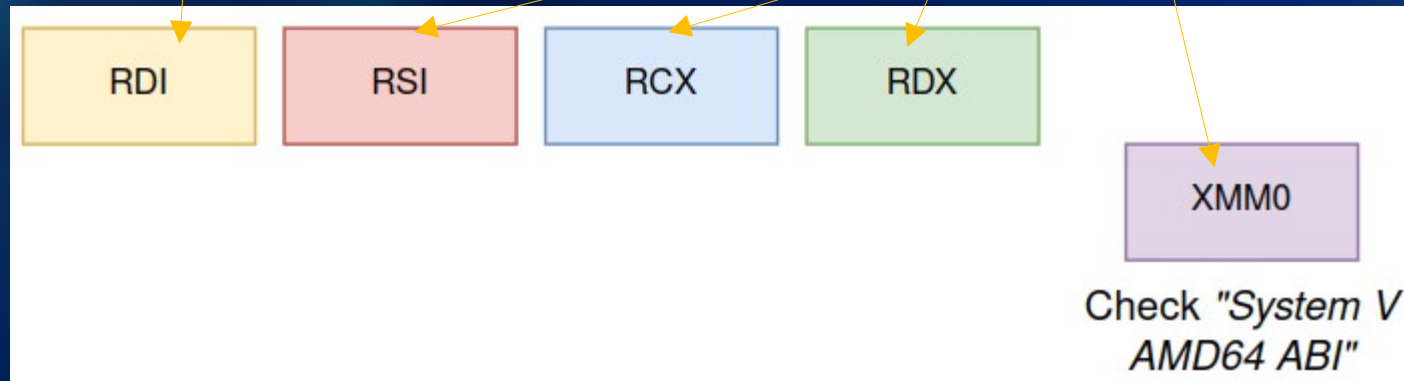
- ▶ Reference: [man 3 printf](#) and [man 3 scanf](#)

The overall syntax of a conversion specification is:

```
%[$][flags][width][.precision][length modifier]conversion
```

- ▶ Example (x86\_64)

```
printf("a=%d x=%lu s1=%s d=%.6lf \n", a, x, s1, d);
```



# Format string at assembly level

```
printf("a=%d x=%lu s1=%s d=%.6lf \n", a, x, s1, d);
```

```
mov dword [a], 3 ; store a
mov eax, 0xaabbccdd ; store x
mov qword [x], rax
lea rax, str.this_is_a_string ; load string address
mov qword [s1], rax
movsd xmm0, qword [0x00002038] ; load double constant
movsd qword [d], xmm0
mov rsi, qword [d] ; store d
mov rcx, qword [s1]
mov rdx, qword [x]
mov eax, dword [a]
movq xmm0, rsi
mov esi, eax
lea rax, str.a_d_x_lu_s1_s_d_.6lf ; load format string
mov rdi, rax ; core arguments
mov eax, 1
call sym.imp.printf ;[1]
```

```
1 #include <stdio.h>
2
3
4 int main(int argc, char *argv[])
5 {
6     int a = 3;
7     unsigned long x = 0xAABBCCDD;
8     char *s1 = "this_is_a_string";
9     double d = 0.0000123;
10
11     printf("a=%d x=%lu s1=%s d=%.6lf \n", a, x, s1, d);
12     return 0;
13 }
```

# Security issue

- ▶ The number of format specifier in the string determine how far into the stack we need to look
- ▶ **Issue:** there is no control on the number of format specifier and the number of argument supplied

```
1 #include <stdio.h>
2
3 int main(int argc, char *argv[])
4 {
5     printf("%d %d %d %d %d\n", 1, 2, 3);
6     return 0;
7 }
```

\$ ./example2  
1 2 3 0 1886076656

# What is happening?

```
1 #include <stdio.h>
2
3 int main(int argc, char *argv[])
4 {
5     printf("%d %d %d %d %d\n", 1, 2, 3);
6     return 0;
7 }
```

Registers sequence for integer arguments:

RDI  
RSI  
RDX  
RCX  
R8  
R9

```
gef> ni
1 2 3 0 -134422800
```

```
0x55555555157 <main+001e>    lea    rax, [rip+0xea6]    # 0x5
0x5555555515e <main+0025>    mov    rdi, rax
0x55555555161 <main+0028>    mov    eax, 0x0
0x55555555166 <main+002d>    call   0x55555555030 <printf@plt>
```

```
$rax : 0x0
$rbx : 0x00007fffffffdb98
$rcx : 0x3
$rdx : 0x2
$rsp : 0x00007fffffffda60
$rbp : 0x00007fffffffda70
$rsi : 0x1
$rdi : 0x000055555556004
$rip : 0x000055555555166
$r8  : 0x0
$r9  : 0x00007ffff7fcdef0
$r10 : 0x00007fffffff790
$r11 : 0x203
$r12 : 0x1
$r13 : 0x0
$r14 : 0x00007ffff7ffd000
$r15 : 0x000055555557dd8
```

**Consequence:**

We can read registers and the stack contents!



# Vulnerability

- ▶ A format string vulnerability is present when the attacker can gain control over the format string supplied to a function

```
char buffer[256];  
printf("Insert your name: ");  
fgets(buffer, 256, stdin);  
buffer[255] = '\\0';  
  
printf(buffer);  
  
return 0;
```

# Example exploitation – Stack leakage

## Normal use

```
> ./example3
Insert your name: hello
hello
```

## Example code

```
1 #include <stdio.h>
2
3 int main(int argc, char *argv[])
4 {
5     char buffer[256];
6     printf("Insert your name: ");
7     fgets(buffer, 256-1, stdin);
8     printf(buffer);
9     return 0;
10 }
```

## Exploitation → read the stack content

```
> python -c 'print("%16lx_"*20)' | ./example3
Insert your name: 557a988dd6b1_ fbad2088_ 1ff_ 557a988dd729_
31255f786c363125_6c3631255f786c36_5f786c3631255f78_31255f786c363125_6c3631255f786c36_5f786c
f786c3631255f78_31255f786c363125_6c3631255f786c36_5f786c3631255f78_31255f786c363125_
```

%16lx → 64 bit value in hex

# What's on the stack

```
> python -c 'print("%16lx_"*20)' | ./example3
Insert your name: 557a988dd6b1_ fbad2088_ 1ff_ 557a988dd729_
31255f786c363125_6c3631255f786c36_5f786c3631255f78_31255f786c363125_6c3631255f786c36_5f786c
f786c3631255f78_31255f786c363125_6c3631255f786c36_5f786c3631255f78_31255f786c363125_
```

The format string itself!

RSI

EDX

RCX

R8

...Then r9 and stack

```
Insert your name: 555555596b1_ fbad2088_ 7fffffff930_0_1_ 7fffffffdb68_ 100000002_5f786c255f786c2
5_5f786c255f786c25_5f786c255f786c25_5f786c255f786c25_5f786c255f786c25_5f786c255f786c25_5f786c25
5f786c25_5f786c255f786c25_5f786c255f786c25_5f786c255f786c25_5f786c255f786c25_5f786c255f786c25_5
f786c255f786c25_5f786c255f786c25_5f786c255f786c25_5f786c255f786c25_5f786c255f786c25_5f786c255f7
86c25_5f786c255f786c25_5f786c255f786c25_5f786c255f786c25_5f786c255f786c25_5f786c255f786c25_5f78
6c255f786c25_5f786c255f786c25_5f786c255f786c25_5f786c255f786c25_5f786c255f786c25_5f786c255f786c
25_5f786c255f786c25_5f786c255f786c25_6c255f786c25_ 7fffffffdb20_ 8827b4bca99a0500_ 7fffffffdae0_ 7f
fff7de0d4a_ 7fffffffda90_ 7fffffffdb68_ 155554040_ 55555555169_ 7fffffffdb68_ 63026182d7579e1c_1_0_7
ffff7fffd000_ 555555557dd8_ 6
```

- Base addr
- Stack addr
- Lib addr (ld.so)
- Lib addr (libc)

0x0000555555554000	0x0000555555555000	0x00000000000001000	r--
0x0000555555555000	0x00005555555556000	0x00000000000001000	r-x
0x00005555555556000	0x00005555555557000	0x00000000000001000	r--
0x00005555555557000	0x00005555555558000	0x00000000000001000	r--
0x00005555555558000	0x00005555555559000	0x00000000000001000	rw-
0x00007ffff7db8000	0x00007ffff7dbb000	0x00000000000003000	rw-
0x00007ffff7dbb000	0x00007ffff7ddf000	0x000000000000024000	r-- /usr/lib/libc.so.6
0x00007ffff7ddf000	0x00007ffff7f43000	0x0000000000000164000	r-x /usr/lib/libc.so.6
0x00007ffff7f43000	0x00007ffff7f91000	0x00000000000004e000	r-- /usr/lib/libc.so.6
0x00007ffff7f91000	0x00007ffff7f95000	0x00000000000004000	r-- /usr/lib/libc.so.6
0x00007ffff7f95000	0x00007ffff7f97000	0x00000000000002000	rw- /usr/lib/libc.so.6
0x00007ffff7f97000	0x00007ffff7fa1000	0x0000000000000a000	rw-
0x00007ffff7fc3000	0x00007ffff7fc7000	0x00000000000004000	r-- [vvar]
0x00007ffff7fc7000	0x00007ffff7fc9000	0x00000000000002000	r-x [vdso]
0x00007ffff7fc9000	0x00007ffff7fca000	0x00000000000001000	r-- /usr/lib/ld-linux-x86-64.so.2
0x00007ffff7fca000	0x00007ffff7ff1000	0x000000000000027000	r-x /usr/lib/ld-linux-x86-64.so.2
0x00007ffff7ff1000	0x00007ffff7ffb000	0x0000000000000a000	r-- /usr/lib/ld-linux-x86-64.so.2
0x00007ffff7ffb000	0x00007ffff7ffd000	0x00000000000002000	r-- /usr/lib/ld-linux-x86-64.so.2
0x00007ffff7ffd000	0x00007ffff7fff000	0x00000000000002000	rw- /usr/lib/ld-linux-x86-64.so.2
0x00007ffff7fff000	0x00007ffff7fff000	0x000000000000022000	rw- [stack]

Executable

# Adding write access (1 of 2)

- ▶ Among the format specifiers, the **%n** works differently: it allow to write inside the corresponding argument (treated as int \*) the number of characters written so far
- ▶ Variants:
  - **%hn** → expects a short \*
  - **%hhn** → expects a char \*
  - **%ln** → expects a long \*

# Adding write access (2 of 2)

- ▶ If we control the format string, we can **forge** one such that:
  - Contains enough format specifiers **to reach the format string itself**
  - Contains the bytes (little endian) of a **valid address**
  - Given **n** the value we would like to write, write **n characters to output**
  - Contains a **%n** (or equivalent) **corresponding to the argument number of the address inserted**
- ▶ This allow an attacker to write a attacker-controlled value to a attacker-controlled address
  - Write-What-Where
- ▶ **Practically:** forge such a format string may require some time to align the address/value correctly



# Optimization: the %m\$ syntax

- ▶ In order to reach the format string in the stack, we need to insert multiple format specifier
  - Too many characters!
- ▶ We can use a special syntax: `%m$f` where `m` is the argument number directly and `f` the format specifier
  - So, if you find the format string after (say) the 123th argument, instead of inserting 122 dummy format specifiers and the `%n`, you just write `%123$n`

# Optimization: write many characters

- ▶ The *value* to write is the result of the number of characters written.
- ▶ But to write a full 64-bit value (say `0xAA03CCDDEEFF1122`) we would need to write **multiple terabytes of data** (!!)
- ▶ To avoid this issue, we can use the *field length* indicator:
  - `%3000c` → write 3000 characters (2999 spaces and a `c` char probably)
- ▶ Normally, it is better to write values one byte at a time (`%hhn`) while putting multiple consecutive addresses in the format string
  - Optionally, ordering the writes such that you only require incremental values may decrease the format string length

# Example

- ▶ Assume you want to write **0xAA03CCDDEEFF1122** at address **0x000000601020**
- ▶ Start from the smallest byte value (the **0x11 = 17**) → **0xAA03CCDDEEFF1122** → address **0x000000601021**  
**%17c%hhn**
- ▶ **0x11** character written, let's proceed with the 2nd smallest – **0x22** – at address **0x000000601020**. We need **0x22-0x11 = 17** chars:  
**%17c%hhn%17c%hhn**
- ▶ Now **0xAA** at address **0x000000601027**. We need **0xAA-0x22 = 126** chars  
**%17c%hhn%17c%hhn%126c%hhn**
- ▶ **0x03** at address **0x000000601026** - **Overflow!** **0x100-0xAA + 3 = 89** chars  
**%17c%hhn%17c%hhn%126c%hhn%89%c%hhn**
- ▶ (and so on ..)
- ▶ Finally, prepend/append all the addresses as bytes in little endian:  
**\x21\x10\x60\x00\x00\x00\x00\x00\x20\x10\x60\x00 ....**

# Pwntools.fmtstr

- ▶ The pwntools suite include a module (fmtstr) to automatically build format strings

- <https://docs.pwntools.com/en/stable/fmtstr.html>

- ▶ Example usage:

```
>>> pwn.fmtstr_payload(5, {0x601020: 0xAABBCCDD}, numwritten=0, write_size='byte')
```

- ▶ Result:

```
%170c%17$hhn%17c%18$hhn%17c%19$hhn%17c%20$  
hhnaaa#\x10`\x00"\x10`\x00!\x10`\x00  
\x10`\x00
```

# Heap Overflow Vulnerabilities

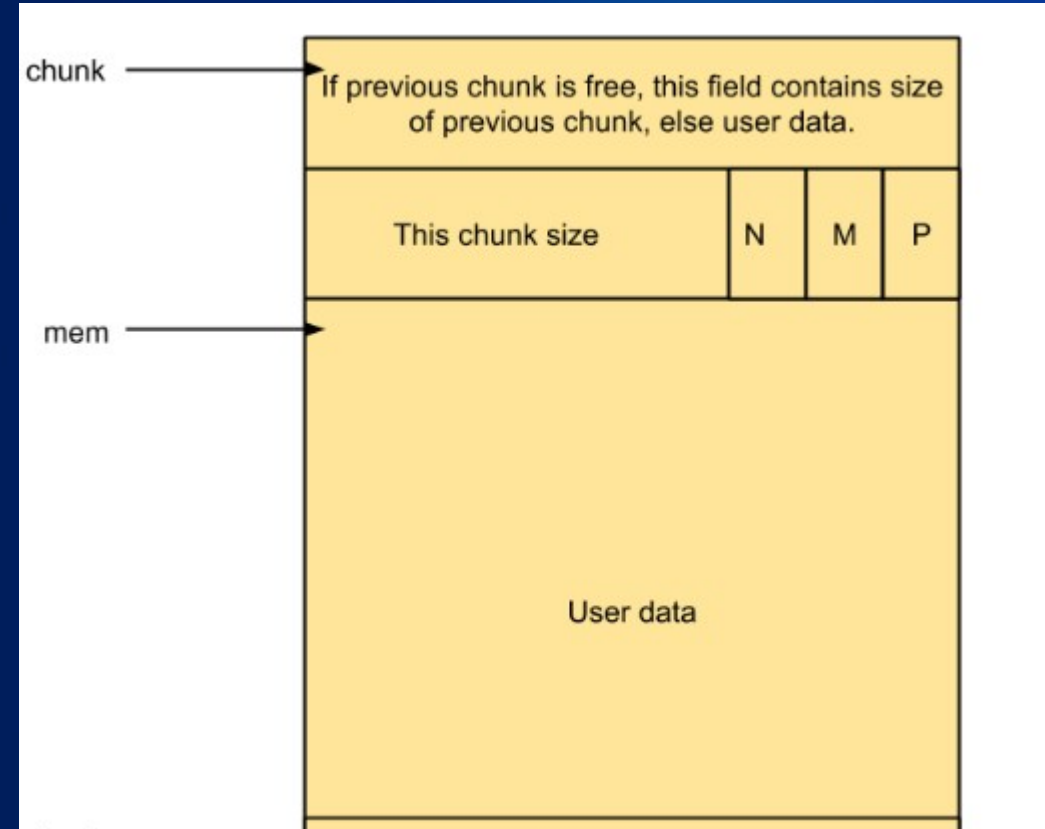


# Heap overflow (3 slides introduction)

- ▶ Heap: conventional name use to indicate the area of memory reserved for dynamic memory allocations
- ▶ The allocation of memory is:
  - Managed by the OS (at memory page level)
  - By libraries (C-library) at user level
- ▶ The component that takes care of memory allocation is the dynamic allocator
  - Many different algorithms/techniques
  - Glibc uses a derivative of the dough-lea allocator

# Heap overflow (3 slides introduction)

- ▶ Memory is allocated with `malloc()` and similar functions. Later, can be de-allocated with `free()`
- ▶ The “piece” of memory assigned to a successful call to `malloc` is called **chunk**
- ▶ In memory, an allocated chunk look like this.
- ▶ When the chunk is de-allocated, it is added to a double-linked list (**bin**) to being reused as soon as required or **coalesced** with other adjacent free chunks.
- ▶ In free chunks, the first two words in the “user data” are used to store the **BK** and **FD** pointers of the double linked list.
- ▶ The **P** bit is used to signal that the previous adjacent chunk is free



# Heap overflow (3 slides introduction)

- ▶ A **Heap overflow** attack consists in overflowing a allocated chunk to overwrite the metadata of the next (adjacent) chunk and to create inconsistencies
- ▶ A successful attack can:
  - Forge non-existing free chunks (later obtainable with malloc() )
  - Alter the size of free/allocated chunks
  - Overlap chunks
  - Force the allocator to give a stack/executable memory area as result of a malloc()
- ▶ Typical attacks
  - Use after free
  - Double Free
  - Fastbin attacks
  - Tcache poisoning
  - Large bin attacks
  - “House of” attacks
- ▶ <https://sploitfun.wordpress.com/2015/02/10/understanding-glibc-malloc/>

# Exploitation on Windows platforms

# Exploitation on Windows – Buffer Overflow

- ▶ The theory of Buffer Overflow is similar
- ▶ The **ABI** and the calling convention in Windows (64-bit) is different:
  - First 4 arguments are passed via different registers
  - Stack should be always aligned to **16 bytes** (2 qwords)
  - Shadow space in the stack
- ▶ Additionally, you cannot count on the C-library. You need to rely on **Win32 API** and/or the functions exported by the **DLL** loaded
  - Normally this is enough since you can use **LoadLibrary+GetProcAddress** to load any DLL and get the pointer to any function



# WIN64 ABI

- ▶ <https://learn.microsoft.com/it-it/cpp/build/x64-software-conventions?view=msvc-170>
- ▶ Register order (for integer arguments):
  - RCX → RDX → R8 → R9
- ▶ Access API:
  - Direct call via Import Address Table (**IAT**)
  - LoadLibrary+GetProcAddress
  - KernelBase DLL access via PEB
- ▶ <https://h0mbre.github.io/Babys-First-Shellcode/>

<https://securitycafe.ro/2015/12/14/introduction-to-windows-shellcode-development-part-2/>

# ROP in Windows

- ▶ Same idea
- ▶ **Keep the stack alignment** is still a requirement/limitation
- ▶ The gadgets required may be difficult to find
  - It is easier to launch another executable than executing a ropchain
  - Alternatives:
    - **Process Hollowing**: debug, stop and map the executable itself, erase/modify, relaunch execution
    - **Process Doppelganging**: similar, but instead of leave the process modified, copy and launch a separate thread/process, rolling back modifications

# Format Strings in Windows

- ▶ Some format specifiers are not available in Windows
  - %n
  - Which Format Specifiers are supported depends also on which C library is used if any (e.g., **CRTDLL.DLL**, **MSVSCRT.DLL** etc.)
- ▶ Still, using format strings for **leaking** is useful
- ▶ Some APIs outside the C libraries uses/accept format strings

# Exercises & Question time